

Essay_4

Wine Quality Prediction with KNN

Authors: Preston O'Connor, Anthony Yasan, Matthew Jacob, Khoa Dao, Nick Wierzbowski

Date: 4/14/2025

Introduction

Our model uses the K nearest neighbor method to see the relation between chemical properties of different wines and the quality ratings of wine experts. The data set is composed of these chemical properties and quality ratings of northern Portuguese 'Vinho Verde' wines, which seems to have been gathered for a research paper similar to ours by Cortez et al. in 2009. The relationship between the chemical properties of wine and the "subjective" flavor, encompassing taste and smell, actually seems to be pretty complex. According to Shapin, who wrote a paper detailing how the American wine tasting language and process was born out of contributions from wine experts, producers and chemists, the flavor or quality is not only affected by flavorants and odorants, which often only have an impact on the smell, but also much less reliable prior experience and expectations. In fact the current language popularized by people like Maynard Amerine and his associates in the post-war period in itself influences how we perceive different flavors, and potentially thus the impact of the chemical flavorants on the quality. All of that is to say that the taste of wine is influenced not just by measurable chemicals but also social factors, especially when considering the professional wine tasting industry and a common language used by them.

The packages we used are `class`, which has functions for categorization useful for KNN implementation.

As for our results, the model seems to have been fairly successful at predicting wine quality when it was mid range, and more common. For the lower and higher ends of quality ratings the model performed worse, which can both be attributable to lower amounts of data for those

categories and perhaps increasing deviation among those categories the further they get from an 'average' wine.

```
#libraries given in Step 1 of the slides
# install.packages("class") # For KNN
# install.packages("tidyverse") # For visualization
# install.packages("corrplot") # correlation matrix visualization
# install.packages("ggplot2")# used to view correltation matrix
# install.packages("leaps") #for best subset sum
# install.packages("pROC") # for the ROC analaysys
library(class) # for KNN Implementation
library(tidyverse) # for visualization
```

```
-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
v dplyr      1.1.4      v readr      2.1.5
v forcats    1.0.0      v stringr    1.5.1
v ggplot2     3.5.1      v tibble     3.2.1
v lubridate  1.9.4      v tidyr      1.3.1
v purrr      1.0.4
```

```
-- Conflicts ----- tidyverse_conflicts() --
```

```
x dplyr::filter() masks stats::filter()
```

```
x dplyr::lag()     masks stats::lag()
```

```
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become
```

```
library(corrplot)
```

corrplot 0.95 loaded

```
library(ggplot2)
library(caret) # for KNN
```

Loading required package: lattice

Attaching package: 'caret'

The following object is masked from 'package:purrr':

lift

```
library(leaps)
library(pROC)
```

Type 'citation("pROC")' for a citation.

Attaching package: 'pROC'

The following objects are masked from 'package:stats':

cov, smooth, var

```
#loading the data set
data <- read.csv("winequality-red.csv")
#names(wine)
summary(data)
```

fixed.acidity	volatile.acidity	citric.acid	residual.sugar
Min. : 4.60	Min. : 0.1200	Min. : 0.000	Min. : 0.900
1st Qu.: 7.10	1st Qu.: 0.3900	1st Qu.: 0.090	1st Qu.: 1.900
Median : 7.90	Median : 0.5200	Median : 0.260	Median : 2.200
Mean : 8.32	Mean : 0.5278	Mean : 0.271	Mean : 2.539
3rd Qu.: 9.20	3rd Qu.: 0.6400	3rd Qu.: 0.420	3rd Qu.: 2.600
Max. : 15.90	Max. : 1.5800	Max. : 1.000	Max. : 15.500
chlorides	free.sulfur.dioxide	total.sulfur.dioxide	density
Min. : 0.01200	Min. : 1.00	Min. : 6.00	Min. : 0.9901
1st Qu.: 0.07000	1st Qu.: 7.00	1st Qu.: 22.00	1st Qu.: 0.9956
Median : 0.07900	Median : 14.00	Median : 38.00	Median : 0.9968
Mean : 0.08747	Mean : 15.87	Mean : 46.47	Mean : 0.9967
3rd Qu.: 0.09000	3rd Qu.: 21.00	3rd Qu.: 62.00	3rd Qu.: 0.9978
Max. : 0.61100	Max. : 72.00	Max. : 289.00	Max. : 1.0037
pH	sulphates	alcohol	quality
Min. : 2.740	Min. : 0.3300	Min. : 8.40	Min. : 3.000
1st Qu.: 3.210	1st Qu.: 0.5500	1st Qu.: 9.50	1st Qu.: 5.000
Median : 3.310	Median : 0.6200	Median : 10.20	Median : 6.000
Mean : 3.311	Mean : 0.6581	Mean : 10.42	Mean : 5.636
3rd Qu.: 3.400	3rd Qu.: 0.7300	3rd Qu.: 11.10	3rd Qu.: 6.000
Max. : 4.010	Max. : 2.0000	Max. : 14.90	Max. : 8.000

Data Description

Data source

The dataset used in this project is the Red Wine Quality dataset from the UCI Machine Learning Repository, also available on Kaggle (<https://www.kaggle.com/datasets/uciml/red-wine-quality-cortez-et-al-2009>). This dataset was compiled by Paulo Cortez et al. (2009) and is related to red Vinho Verde wine samples from the north of Portugal.

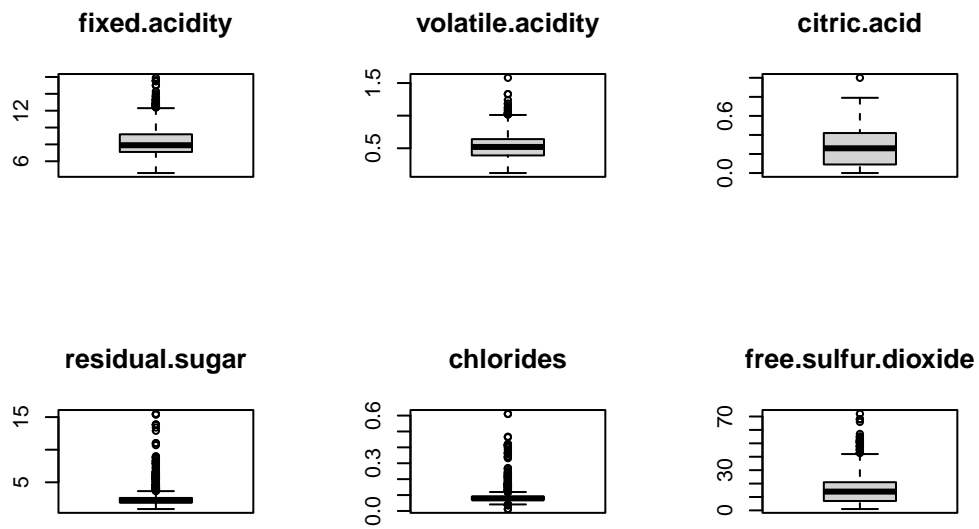
Data structure and Variables

The dataset contains 1599 observations (rows) and 12 variables (11 numeric input variables and 1 output variable: quality). The variables represent physicochemical attributes of the wine samples:

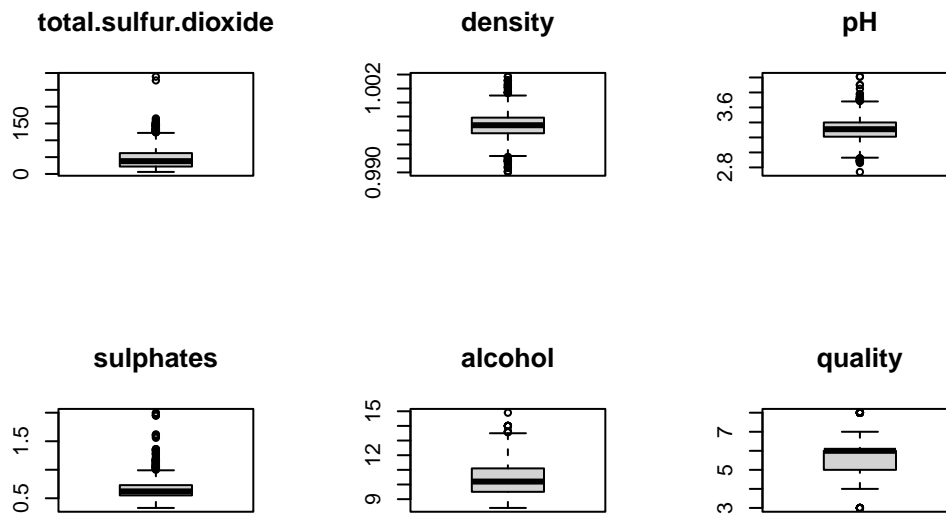
- **fixed acidity:** Refers primarily to tartaric acid, which is one of the main acids found naturally in grapes.
- **volatile acidity:** Measures the amount of acetic acid (vinegar) in the wine.
- **citric acid:** A natural acid found in small quantities in wine that can add freshness and flavor.
- **residual sugar:** Represents the amount of sugar remaining after fermentation.
- **chlorides:** Indicates the amount of salt in the wine, which can affect the taste and preservation.
- **free sulfur dioxide:** Refers to the part of sulfur dioxide (SO_2) that is not bound to other molecules and is available to act as an antioxidant and antimicrobial agent.
- **total sulfur dioxide:** Includes both free and bound forms of SO_2 .
- **density:** The density of wine, which is influenced by the sugar and alcohol content.
- **pH:** Indicates how acidic or basic the wine is.
- **sulphates:** Sulfate compounds can contribute to the wine's flavor and preservation.
- **alcohol:** The percentage of ethanol by volume in the wine.
- **quality:** This is the response variable, a sensory score assigned by professional tasters ranging from 0 to 10. It reflects the overall quality of the wine sample based on taste, aroma, and balance.

Data cleaning

```
# Boxplots to detect outliers
par(mfrow = c(2, 3))
for (i in 1:6) {
  boxplot(data[[i]], main = names(data)[i])
}
```



```
# Plot remaining 6 variables
par(mfrow = c(2, 3))
for (i in 7:12) {
  boxplot(data[[i]], main = names(data)[i])
}
```



```
par(mfrow = c(1, 1)) # Reset

# Function to remove outliers beyond 1.5 * IQR
remove_outliers_IQR <- function(df, column) {
  #df <- df %>% filter(!is.na(df[[column]])) # Remove missing data
  df <- df %>% filter(!is.na(!sym(column)))
  Q1 <- quantile(df[[column]], 0.25, na.rm = T)
  Q3 <- quantile(df[[column]], 0.75, na.rm = T)
  IQR <- Q3 - Q1
  lower_bound <- Q1 - 1.5 * IQR
  upper_bound <- Q3 + 1.5 * IQR
  df %>% filter(df[[column]] >= lower_bound & df[[column]] <= upper_bound)
}

# Remove outliers for all relevant columns
data_cleaned <- data
columns <- c("fixed.acidity", "volatile.acidity", "citric.acid",
             "residual.sugar", "chlorides", "free.sulfur.dioxide",
             "total.sulfur.dioxide", "density", "pH",
             "sulphates", "alcohol")

data_cleaned <- data
```

```
for (col in columns) {
  data_cleaned <- remove_outliers_IQR(data_cleaned, col)
}
```

```
# Summary of the cleaned data
summary(data_cleaned)
```

fixed.acidity	volatile.acidity	citric.acid	residual.sugar
Min. : 5.100	Min. :0.1200	Min. :0.000	Min. :1.200
1st Qu.: 7.100	1st Qu.:0.3900	1st Qu.:0.080	1st Qu.:1.900
Median : 7.800	Median :0.5200	Median :0.240	Median :2.100
Mean : 8.147	Mean :0.5222	Mean :0.246	Mean :2.181
3rd Qu.: 9.000	3rd Qu.:0.6300	3rd Qu.:0.390	3rd Qu.:2.400
Max. :12.300	Max. :1.0050	Max. :0.730	Max. :3.600
chlorides	free.sulfur.dioxide	total.sulfur.dioxide	density
Min. :0.0420	Min. : 1.00	Min. : 6.00	Min. :0.9926
1st Qu.:0.0690	1st Qu.: 8.00	1st Qu.: 22.00	1st Qu.:0.9955
Median :0.0780	Median :13.00	Median : 35.00	Median :0.9966
Mean :0.0783	Mean :14.79	Mean : 40.95	Mean :0.9966
3rd Qu.:0.0870	3rd Qu.:20.00	3rd Qu.: 54.00	3rd Qu.:0.9975
Max. :0.1160	Max. :40.00	Max. :113.00	Max. :1.0004
pH	sulphates	alcohol	quality
Min. :2.980	Min. :0.3300	Min. : 8.70	Min. :3.000
1st Qu.:3.230	1st Qu.:0.5500	1st Qu.: 9.50	1st Qu.:5.000
Median :3.330	Median :0.6100	Median :10.10	Median :6.000
Mean :3.325	Mean :0.6284	Mean :10.35	Mean :5.637
3rd Qu.:3.400	3rd Qu.:0.7000	3rd Qu.:11.00	3rd Qu.:6.000
Max. :3.680	Max. :0.9400	Max. :13.00	Max. :8.000

```
# New size of the cleaned dataset
nrow(data_cleaned)
```

```
[1] 1135
```

The dataset contains 1135 observations (rows) after removing outliers.

Normalize Data

```
normalize <- function(x) { (x - min(x)) / (max(x) - min(x)) }

data_norm <- as.data.frame(lapply(data_cleaned[, 1:11], normalize))
head(data_norm)
```

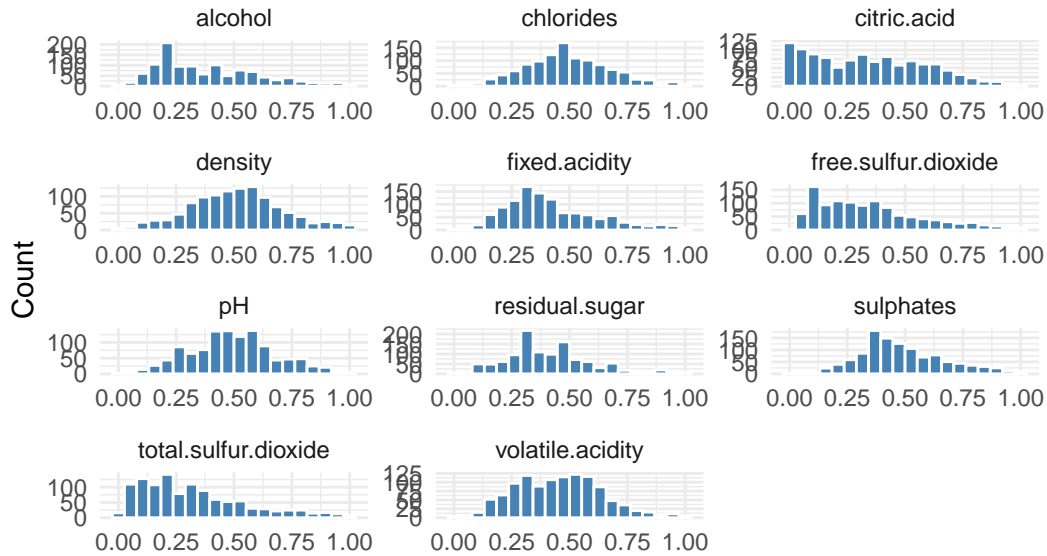
```
fixed.acidity volatile.acidity citric.acid residual.sugar chlorides
1      0.3194444      0.6553672  0.00000000      0.2916667 0.4594595
2      0.3750000      0.8587571  0.00000000      0.5833333 0.7567568
3      0.3750000      0.7231638  0.05479452      0.4583333 0.6756757
4      0.8472222      0.1807910  0.76712329      0.2916667 0.4459459
5      0.3194444      0.6553672  0.00000000      0.2916667 0.4594595
6      0.3194444      0.6101695  0.00000000      0.2500000 0.4459459
free.sulfur.dioxide total.sulfur.dioxide density      pH sulphates
1      0.2564103      0.2616822 0.6683673 0.7571429 0.3770492
2      0.6153846      0.5700935 0.5408163 0.3142857 0.5737705
3      0.3589744      0.4485981 0.5663265 0.4000000 0.5245902
4      0.4102564      0.5046729 0.6938776 0.2571429 0.4098361
5      0.2564103      0.2616822 0.6683673 0.7571429 0.3770492
6      0.3076923      0.3177570 0.6683673 0.7571429 0.3770492
alcohol
1 0.1627907
2 0.2558140
3 0.2558140
4 0.2558140
5 0.1627907
6 0.1627907
```

Data visualization

```
# Reshape data to long format
wine_long <- pivot_longer(data_norm, cols = 1:11, names_to = "variable",
  ↪ values_to = "value")

# Create faceted histogram plot
ggplot(wine_long, aes(x = value)) +
  geom_histogram(bins = 20, fill = "steelblue", color = "white") +
  facet_wrap(~variable, scales = "free", ncol = 3) +
  theme_minimal() +
  labs(title = "Histograms of Wine Features", x = "", y = "Count")
```


Histograms of Wine Features



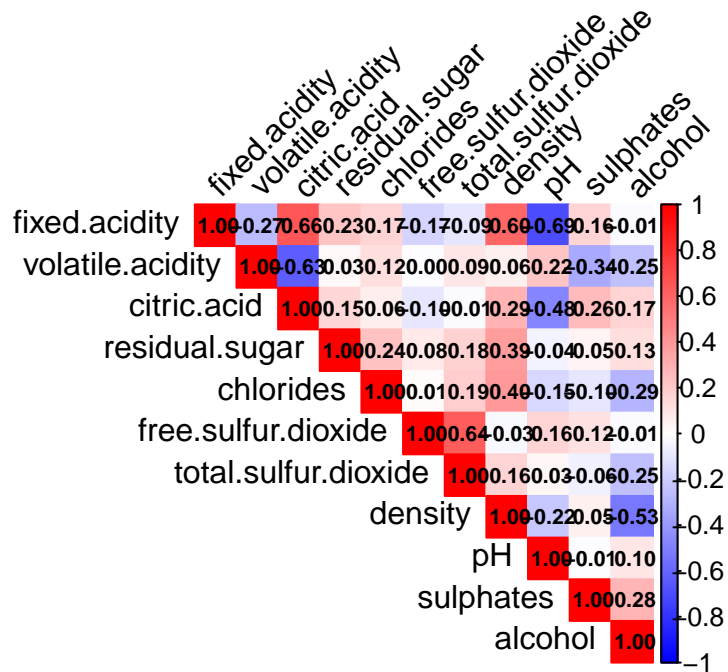
Correlation matrix

```
# Compute correlation matrix for numeric input variables
cor_matrix <- cor(data_norm[, 1:11]) # Exclude 'quality' if you're only
  ↳ interested in inputs
print(round(cor_matrix, 2))          # Round for readability
```

	fixed.acidity	volatile.acidity	citric.acid	residual.sugar
fixed.acidity	1.00	-0.27	0.66	0.23
volatile.acidity	-0.27	1.00	-0.63	0.03
citric.acid	0.66	-0.63	1.00	0.15
residual.sugar	0.23	0.03	0.15	1.00
chlorides	0.17	0.12	0.06	0.24
free.sulfur.dioxide	-0.17	0.00	-0.10	0.08
total.sulfur.dioxide	-0.09	0.09	-0.01	0.18
density	0.60	0.06	0.29	0.39
pH	-0.69	0.22	-0.48	-0.04
sulphates	0.16	-0.34	0.26	0.05
alcohol	-0.01	-0.25	0.17	0.13
	chlorides	free.sulfur.dioxide	total.sulfur.dioxide	density

fixed.acidity	0.17	-0.17	-0.09	0.60
volatile.acidity	0.12	0.00	0.09	0.06
citric.acid	0.06	-0.10	-0.01	0.29
residual.sugar	0.24	0.08	0.18	0.39
chlorides	1.00	0.01	0.19	0.40
free.sulfur.dioxide	0.01	1.00	0.64	-0.03
total.sulfur.dioxide	0.19	0.64	1.00	0.16
density	0.40	-0.03	0.16	1.00
pH	-0.15	0.16	0.03	-0.22
sulphates	-0.10	0.12	-0.06	0.05
alcohol	-0.29	-0.01	-0.25	-0.53
	pH	sulphates	alcohol	
fixed.acidity	-0.69	0.16	-0.01	
volatile.acidity	0.22	-0.34	-0.25	
citric.acid	-0.48	0.26	0.17	
residual.sugar	-0.04	0.05	0.13	
chlorides	-0.15	-0.10	-0.29	
free.sulfur.dioxide	0.16	0.12	-0.01	
total.sulfur.dioxide	0.03	-0.06	-0.25	
density	-0.22	0.05	-0.53	
pH	1.00	-0.01	0.10	
sulphates	-0.01	1.00	0.28	
alcohol	0.10	0.28	1.00	

```
corrplot(cor_matrix, method = "color", type = "upper",
         tl.col = "black", tl.srt = 45,
         addCoef.col = "black", number.cex = 0.7,
         col = colorRampPalette(c("blue", "white", "red"))(200))
```



Analysis

Best Subset

```
data_combined <- cbind(data_norm, quality = data_cleaned$quality)
nrow(data_norm)
```

```
[1] 1135
```

```
# due to 2 instances of the value of quality 3 remove the data
data_combined <- data_combined %>%
  filter(quality !=3)
# remove the data with the two points of classification as three
```

```
nrow(data_combined)
```

```
[1] 1133
```

Analysis

Train and Test Split with the optimized K

- We used an 80:20 ratio for our training and test sets to provide the model with a sizable amount of data to train on, while still maintaining a good mechanism for evaluating its performance. This yielded a training set of 906 and a testing set of 227. Using all the columns of the data set allows the KNN model to consider every feature that could potentially have an impact on wine quality, resulting in more accurate and informed predictions. Using the full set of variables, we are not in danger of leaving out any important information.

```
# First Implement the Seed and split the Training and Testing Data
set.seed(12)
#selected_cols <- c(2,10,11) # this received 61.25% with k =20
#selected_cols <- c(2,3,4,5, 7, 8, 9,10,11) # this received 61.87% with k =
↪ 31
index <- sample(1:nrow(data_combined), 0.8 * nrow(data_combined))

train_data <- data_combined[index, 1:11] # Features: columns 1 to 11
test_data <- data_combined[-index, 1:11]

cat("Number of training samples:", nrow(train_data), "\n")
```

Number of training samples: 906

```
cat("Number of testing samples:", nrow(test_data), "\n")
```

Number of testing samples: 227

```
train_labels <- data_combined[index, 12] # Target: column 12
test_labels <- data_combined[-index, 12]

# The numbers are treated as categorical not to calculate the average from
train_labels <- as.factor(train_labels)
test_labels <- as.factor(test_labels)

k_values <- 2:60

# Train KNN with k = 2-60
```

```

accuracy_list <- c()

for (k in k_values) {
  knn_pred <- knn(train = train_data, test = test_data, cl = train_labels, k
  ↪ = k)
  acc <- mean(knn_pred == test_labels)
  accuracy_list <- c(accuracy_list, acc)
}

# Find the best k
best_k <- k_values[which.max(accuracy_list)]
best_acc <- max(accuracy_list)

cat("The Best k:", best_k, "\n")

```

The Best k: 6

```

cat("The Best Accuracy:", round(best_acc, 4), "\n")

```

The Best Accuracy: 0.6211

```

# Final model with best k
final_knn <- knn(train = train_data, test = test_data, cl = train_labels, k =
  ↪ best_k)

```

Model Evaluation


The K is optimized

```

set.seed(13)
k_values <- 1:60
accuracies <- sapply(k_values, function(k) {
  pred <- knn(train_data, test_data, cl = train_labels, k = k)
  mean(pred == test_labels)
})
# Plot accuracy vs. K
plot(k_values, accuracies, type = "b", col = "blue", pch = 19,
  xlab = "K", ylab = "Accuracy", main = "Optimal K Selection")

```

Optimal K Selection




- The plot illustrates the accuracy of the KNN model for different values of K, with optimal accuracy at low values of K, namely in the range of K equal to 5. As K increases, the model's accuracy has a tendency to fall, meaning larger neighborhoods reduce the model's ability to make correct predictions. Therefore, a lower value of K is more appropriate for this data set, sacrificing sensitivity to local trends for overall performance.

The Confusion Matrix

```
conf_matrix <- table(Predicted = final_knn, Actual = test_labels)

conf_matrix_df <- as.data.frame(as.table(conf_matrix))

ggplot(data = conf_matrix_df, aes(x = Predicted, y = Actual, fill = Freq)) +
  geom_tile() +
  geom_text(aes(label = Freq), color = "black", size = 6) +
  scale_fill_gradient(low = "white", high = "blue") +
  theme_minimal() +
  labs(x = "Predicted", y = "Actual", fill = "Frequency") +
  theme(axis.text.x = element_text(hjust = 1))
```



- Based on the confusion matrix, the KNN model is moderately accurate with the majority of the correct predictions clustered around classes 5 and 6. Class 6 has the maximum true positive value (61), indicating that the model is most confident in classifying this class correctly; however, there is a large misclassification between neighboring classes: a majority of class 5 wines are forecast as class 6 (64), and the class 6 wines are largely forecast as 5 (24) or 7 (6). Class 7 forecasts are everywhere, and the model performs very badly with classes 4 and 8, correctly classifying none within these groups. This means that the model is most powerful at detecting common, central classes, but less sensitive to less common or edge-case classes. Overall, the model gives general trends regarding wine quality, but may be assisted by algorithms that provide improved performance on underrepresented classes.

Multiple of Class Precision, Recall, and F1 Score

```
conf_matrix <- confusionMatrix(final_knn, test_labels)
# print(conf_matrix) # for debugging purposes

cm_by_class <- conf_matrix$byClass

if (is.null(dim(cm_by_class))) {
  precision <- cm_by_class["Pos Pred Value"]
}
```

```

recall <- cm_by_class["Sensitivity"]
f1 <- 2 * (precision * recall) / (precision + recall)
metrics_df <- data.frame(
  Class = levels(test_labels),
  Precision = round(precision, 3),
  Recall = round(recall, 3),
  F1_Score = round(f1, 3)
)
} else {
precision <- cm_by_class[, "Pos Pred Value"]
recall <- cm_by_class[, "Sensitivity"]
f1 <- 2 * (precision * recall) / (precision + recall)
metrics_df <- data.frame(
  Class = rownames(cm_by_class),
  Precision = round(precision, 3),
  Recall = round(recall, 3),
  F1_Score = round(f1, 3)
)
}

print(metrics_df)

```

	Class	Precision	Recall	F1_Score
Class: 4	Class: 4	0.000	0.000	NaN
Class: 5	Class: 5	0.634	0.681	0.656
Class: 6	Class: 6	0.551	0.648	0.596
Class: 7	Class: 7	0.706	0.353	0.471
Class: 8	Class: 8	NaN	0.000	NaN

- Class 4: The model did not predict any wines as quality 4. Precision and recall are both 0, indicating that it completely missed this class.
- Class 5: For class 5, the model is most accurate, with 65.3% precision and 68.1% recall, meaning it correctly identifies most wines with this quality.
- Class 6: Class 6 predictions are good with precision 56.5% and recall 67%, indicating moderate accuracy and good coverage.
- Class 7: The model performs badly on class 7 with only 50% accuracy and 29.4% recall, thus missing most of the true class 7 wines.
- Class 8: The model doesn't predict class 8 at all. Both precision and recall are 0, which means it completely fails to predict this class.
- delete this point is so i know i submitted the right draft on 4/11

Conclusion & Summary

In this project, we used the K-Nearest Neighbors (KNN) algorithm to predict the quality of Portuguese ‘Vinho Verde’ wines based on a range of properties, including acidity, sugar content, and pH. After some data cleaning, normalization, and visualization, we trained and evaluated our model using an 80:20 training/test split. We determined the optimal values for K via different techniques, and then tested our optimized model using a confusion matrix. Our model achieved reasonable accuracy in predicting mid-range wine qualities, especially for the more common quality ratings (classes 5 and 6). However, the model faced challenges with predicting extreme quality classes (classes 4, 7, and 8). This may be as a result of there being less representation for particularly low quality and high quality wines. Overall, the model is able to accurately suggest certain trends regarding wine quality, particularly for mid-range wines, but it needs bolstering via algorithms that can better incorporate underrepresented classes into the model. This model could also be improved by including additional variables that account for social, environmental, and historical factors that may impact wine quality.

Sources

- Shapin, Steven. “A Taste of Science: Making the Subjective Objective in the California Wine World.” *Social Studies of Science* 46, no. 3 (2016): 436–60. <http://www.jstor.org/stable/26099849>.
- P. Cortez, A. Cerdeira, F. Almeida, T. Matos and J. Reis. Modeling wine preferences by data mining from physicochemical properties. In *Decision Support Systems*, Elsevier, 47(4):547-553, 2009.
- Choudhury, P. (2020, June 26). Best way to learn KNN algorithm using R programming. Analytics Vidhya. <https://www.analyticsvidhya.com/blog/2015/08/learning-concept-knn-algorithms-programming/>